Analysis receipt

Every AI analysis makes compromises. Document yours.

A one-page record that travels with any qualitative analysis, done by people, by AI, or both. It shows what the analysis read, what it set aside, whether anyone checked, and what each finding rests on. It does not certify the work. It lets others judge it.

Four parts, each with an owner

1 Question

What was asked, and what decision it serves.

Whoever set the question

2 Readers

What was read, by whom or by what, and what was never read.

Whoever ran the reading

3 Leftovers

What was set aside, and whether anyone checked it.

Whoever checks the leftovers

4 Claims

What each finding rests on, and the cases against it.

Whoever reports the findings

Five steps

1. **Make a copy before you start.** Write the research question in, word for word, before any reading begins.
2. **Fill it in while you work.** A record made afterwards is a memory. Note each reader as you use it, including any AI tool, its version and the date. Write counts and reasons, not adjectives. Blank reads as "not checked".
3. **Check the leftovers.** The one step that adds new evidence. See the box.
4. **Look for the rare views.** A random sample rarely finds what one or two people said. Ask a second reader, a person or an AI, to find the people the question does not fit. Or read the passages your codes fit worst.
5. **Attach it to what you share.** Put a short version on the last page or slide, and link the full receipt. Each owner signs their part.

Checking the leftovers

A **leftover** is a passage no finding used, such as an interview answer no theme quotes.

1. List the passages no finding used.
2. Pick some at random.
3. Read each one. Does it fit an existing finding, hold a new idea, or hold nothing relevant?
4. Add any new idea to the findings, and say so on the receipt.

READ, NOTHING NEW	MAY STILL HOLD A NEW IDEA
20 leftovers	up to about 14%
30 leftovers	up to about 10%
60 leftovers	up to about 5%
100 leftovers	up to about 3%

At 95% confidence. Pick the size the decision can live with.

Words used here

Reader — Anything that turned material into findings. A person, an AI model, a tool, or a pipeline step.

Passage — The smallest unit that carries one idea, such as one interview answer or one survey response.

Remember. A receipt is not a certificate. A random sample finds common misses, not rare ones. No reading is ground truth, including yours. Agreement between readers is not coverage. The leftover check and the claim checks are what keep it from becoming a ritual.

Template on the next page. Filled example on page 3.

sashamitts.github.io/analysis-receipt · CC BY 4.0 · Sasha Mitts

Analysis receipt

Fill in while you work. Blank reads as "not checked".
If a line does not apply, write "Not applicable" and why.

Study or project

Analysis owner

Dates of the analysis

Receipt version

QUESTION

Whoever set the question

Question and frame. Word for word as given to every reader. Lens used. What the frame leaves out by design.

What it means for the decision. Ship / ship with a stated limit / read more / no decision rests on it.

READERS

Whoever ran the reading

What was collected. Count, kind, source, how selected.

What was read, and by which readers. People, models and tools. Note sampling, truncation, chunking, retrieval, cut-offs. If a tool cannot show what it read, say so.

Reader	Person, model or tool (version, date)	Given	Actually read	Brief

What was never read, and why. Counts and reasons.

LEFTOVERS

Whoever checks what was set aside

What counts as a leftover. The unit (turn, response, note) and the rule (for example, no finding quotes it).

What was sampled, and what it found. The one line that produces new evidence. Zero new in 60 random leftovers: probably under 5%. It bounds common misses, not rare ones.

Total leftovers	Sampled	Seed	Read by	Existing	New	Nothing	Upper bound (95%)

Targeted searches for rare or misfitting material. Briefs or rankings used, and what they found.

What was not checked. Named plainly.

CLAIMS

Whoever reports the findings

What each claim rests on. Support counted from the records, and of how many. Cases against, and what reading them showed.

Claim	Support (n of N)	Cases against	What they showed

Claims that did not hold. Proposed by any reader, checked, dropped or narrowed, and why.

Not a certificate. A record of what the analysis rests on, and what it did not look at.

Analysis receipt v1.0 · CC BY 4.0 · Sasha Mitts

Question signed

Readers signed

Leftovers signed

Claims signed

Analysis receipt, filled example

Learners East 2026 demonstration
Owner Sasha Mitts · v1 · 1 to 8 October 2026

Twelve interviews from the public Anthropic Interviewer dataset, read by a person and by four AI setups, for the talk "Finding What AI Analysis Misses". It is small on purpose. Read it as an example of the method, not as a finding about AI tools in general.

	WHAT IT RESTS ON	WHAT IT LEAVES OPEN
QUESTION	• "How do these professionals decide what to hand to AI, what to keep for themselves, and when to check its work?" • Given word for word to every reader	• No framework applied • Anything beyond handing work over, keeping it or checking it, although several people talked about their jobs more widely • A demonstration. No decision rests on it.
READERS	• 12 interviews drawn at random from 1,000. 16,963 words, 138 participant answers. • R. My own reading, blind, before seeing any AI output • A. One AI model, one pass. A×4. The same model and prompt, four passes. • B. Three AI providers. C. Four separate targeted searches. • Model names, settings and prompts are in the run log	• The other 988 interviews • Each AI call read all 12 interviews at once, in the same order. That suits patterns across people and may cost detail within each person. Reading one interview at a time was not tested. • Temperature was left at each provider's default
LEFTOVERS	• A leftover is an answer no AI finding quotes. 45 of 138. • 20 read at random by me. 8 fit existing ideas, 0 held a new idea, 12 held nothing relevant. • 3 of the 8 were the worry about being displaced, which no AI reader found • The "where the question does not fit" search found 1 person, already in my reading	• Up to about 14% of leftovers may still hold a new idea (95%) • 25 leftovers unread • One human reference reader, and no second
CLAIMS	• One AI pass found 19 of the 34 ideas in my reading (56%). All AI readers together found 27 (79%). • Every idea held by five or more people was found. All 7 misses were held by four or fewer. • All 526 AI quotes were found word for word in the transcripts, checked by code	• 3 ideas only the AI proposed. 2 held partly. 1 was already in my reading. • Claude matched AI findings to my ideas. I checked 20 at random, and all 20 held.

Question · Sasha Mitts Readers · Sasha Mitts Leftovers · Sasha Mitts Claims · Sasha Mitts

Full version with every line: github.com/sashamitts/analysis-receipt, examples folder

CC BY 4.0 · Sasha Mitts